

DSCI 551 Review

Jakob Thoms

2022-10-04

Lecture 1

Basic probability concepts:

In general, the probability of an event A occurring is denoted as $P(A)$ and is defined as

$$P(A) = \frac{\text{Number of times event } A \text{ is observed}}{\text{Total number of events observed}}$$

as the number of events goes to infinity.

- We heavily rely on the “frequency of events” to make estimations of specific parameters of interest in a population or system.
- This is basically the foundation of a frequentist approach: relying on the frequency (or “number”!) of events to estimate your parameters of interest.

Law of total probability: When partitioning the sample space (the set of all possible events), the sum of the probabilities of each event should be one.

$$\sum_{E \in \Omega} P(E) = 1.$$

- In general, for a given event A , the law implies that

$$1 = P(A) + P(A^c).$$

Inclusion-exclusion principle:

$$P(A \cup B) = P(A) + P(B) - P(A \cap B),$$

$$P(A \cup B \cup C) = P(A) + P(B) + P(C) - P(A \cap B) - P(B \cap C) - P(A \cap C) + P(A \cap B \cap C),$$

etc.

Odds: are quite helpful in comparing the probability of two events.

$$o = \frac{p}{1-p},$$

where p is the probability of an event.

- This implies

$$p = \frac{o}{o+1}.$$

Measures of central tendency and uncertainty:

Central tendency: a measure denoting a “typical” value in a random variable.

Uncertainty: a measure of how “spread” a random variable is

- Called **parameters** when it comes to a population
- Are estimated via **sample statistics**

Mode: the outcome having the highest probability (discrete) or highest probability density (continuous)

Entropy: a measure of uncertainty defined by

$$H(X) = \sum_x P(X = x) \ln \left(\frac{1}{P(X = x)} \right)$$

or

$$H(X) = \int_x f_X(x) \ln \left(\frac{1}{f_X(x)} \right) dx.$$

- Always non-negative in the discrete case
- $H(X) = 0 \iff X$ is constant in the discrete case.

Expectation:

$$\mathbb{E}(X) = \sum_x x \cdot P(X = x).$$

or

$$\mathbb{E}(X) = \int_x x \cdot f_X(x)$$

- Can usually be estimated via the **sample mean**

$$\bar{X} = \frac{1}{n} \sum_{i=1}^n X_i$$

Variance:

$$\text{Var}(X) = \mathbb{E}\{[X - \mathbb{E}(X)]^2\}.$$

$$\implies \text{Var}(X) = \mathbb{E}(X^2) - [\mathbb{E}(X)]^2.$$

- the variance is an expectation (specifically, the squared deviation from the mean)
- can usually be estimated via the **sample variance**

$$S^2 = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})^2$$

- Always non-negative, and $\text{Var}(X) = 0 \iff X$ is constant

Standard deviation: The square root of the variance,

$$\sigma_X = \sqrt{\text{Var}(X)}.$$

Lecture 2

- To maximize entropy, you need equal probabilities for all the outcomes in the sample space. This indicates we have a uniform uncertainty over the whole range of possible outcomes.
- Helpful univariate distribution guide: <http://www.math.wm.edu/~leemis/chart/UDR/UDR.html>

Binomial distribution:

$$X \sim \text{Binomial}(n, \pi)$$

- X is the number of successes in n trials in which each trial has probability π of success, independent of all other trials.
- PMF:

$$P(X = x \mid n, \pi) = \binom{n}{x} \pi^x (1 - \pi)^{n-x} \quad \text{for } x = 0, 1, \dots, n.$$

- Expected value:

$$\mathbb{E}(X) = n\pi$$

- Variance:

$$\text{Var}(X) = n\pi(1 - \pi)$$

Families and Parameters:

- We refer to the entire set of Binomial probability distributions as the **Binomial family of distributions**.
- Specifying a value for both π and n results in a unique Binomial distribution.
- Since π and n fully specify a Binomial distribution, we call them **parameters** of the Binomial family, and we call the Binomial family a **parametric family** of distributions.
- There are other ways we can specify the distribution. For instance, specifying the mean and variance is enough to identify a Binomial distribution.
- Exactly which variables we decide to use to identify a distribution within a family is called the family's parameterization.
- The parameterization you use in practice will depend on the information you can more easily obtain

Geometric distribution:

$$X \sim \text{Geometric}(\pi)$$

X is the number of trials **before** experiencing a success, where each trial has probability π of success, independent of all other trials.

- PMF:

$$P(X = x \mid \pi) = \pi(1 - \pi)^x \quad \text{for } x = 0, 1, \dots$$

- Since there is only one parameter, this means that if you know the mean, you also know the variance!
- Expected value:

$$\mathbb{E}(X) = \frac{1 - \pi}{\pi}$$

- Variance:

$$\text{Var}(X) = \frac{1 - \pi}{\pi^2}$$

Negative Binomial Distribution:

$$X \sim \text{Negative Binomial}(k, \pi)$$

- X is the number of failed trials before experiencing k successes, where each trial has probability π of success, independent of all other trials. - PMF:

$$P(X = x | k, \pi) = \binom{k-1+x}{x} \pi^k (1-\pi)^x \quad \text{for } x = 0, 1, \dots$$

- The Geometric family results with $k = 1$.
- Expected value:

$$\mathbb{E}(X) = \frac{k(1-\pi)}{\pi}.$$

- Variance:

$$\text{Var}(X) = \frac{k(1-\pi)}{\pi^2}.$$

Poisson Distribution:

$$X \sim \text{Poisson}(\lambda)$$

- X is number of events occurring in a fixed interval of time or space, assuming that these events occur with a known constant mean rate (e.g. 3 events per minute or 5 events per meter) and independently of the time since the last event
- PMF

$$P(X = x | \lambda) = \frac{\lambda^x \exp(-\lambda)}{x!} \quad \text{for } x = 0, 1, \dots$$

- Expected value:

$$\mathbb{E}(X) = \lambda.$$

- Variance:

$$\text{Var}(X) = \lambda.$$

Bernoulli Distribution:

$$X \sim \text{Bernoulli}(\pi)$$

- X is equal to one with probability π and equal to zero with probability $1 - \pi$.
- Basically a weighted coin-flip
- A special case of the Binomial family ($n = 1$)
- PMF:

$$P(X = x | \pi) = \pi^x (1-\pi)^{1-x} \quad \text{for } x = 0, 1.$$

- Expected value:

$$\mathbb{E}(X) = \pi.$$

- Variance:

$$\text{Var}(X) = \pi(1-\pi).$$

Lecture 3

Joint distributions and marginal distributions:

- A **joint distribution** is the distribution of n -tuples of random variables, where $n \geq 2$.
- The distribution of an individual variable is called the **marginal distribution** (sometimes just “marginal” or “margin”).
- The word “marginal” is not really needed when we are talking about a standalone random variable – there is no difference between the “marginal distribution of X ” and the “distribution of X .” Therefore, we just use the word “marginal” to emphasize that the distribution is being considered in isolation from other related variables in the same process or system.
- Going from the initial marginal distributions to the joint distribution is not a straightforward procedure.
- It requires us to understand the dependency structure among the random variables
- If we assume that all the RVs are independent, then we can just multiply the probabilities from the marginal distributions to find the joint distribution
- If you have a joint distribution, then the marginal distribution of each individual variable follows as a consequence
- Just sum up (discrete) or integrate (continuous), and apply the law of total probability:

$$P(A) = \sum_n P(A \cap B_n),$$

or

$$P(A) = \int_y P(A \cap Y = y).$$

Independence:

- X and Y are **independent** if

$$P(X = x \cap Y = y) = P(X = x) \cdot P(Y = y) \quad \forall_{x,y}$$

- Equivalently:

$$P(X = x \mid Y = y) = P(X = x) \quad \forall_{x,y}$$

- In other words: X and Y are independent if knowing something about one of them tells us nothing about the other.

Dependence Measures:

Covariance:

$$\text{Cov}(X, Y) = \mathbb{E}[(X - \mu_X)(Y - \mu_Y)],$$

where $\mu_X = \mathbb{E}(X)$ and $\mu_Y = \mathbb{E}(Y)$.

$$\implies \text{Cov}(X, Y) = \mathbb{E}(XY) - \mathbb{E}(X)\mathbb{E}(Y).$$

- Note that

$$\mathbb{E}(XY) = \mathbb{E}(X)\mathbb{E}(Y) \iff \text{Cov}(X, Y) = 0,$$

- Also if X and Y are independent then $\text{Cov}(X, Y) = 0$.
- But the reverse implication does **not** hold in general!
- Zero covariance simply indicates that there is no **linear** trend

Pearson's correlation:

- $\text{Cov}(X, Y)$ is dependent on the scale of X and Y .
- e.g. $\text{Cov}(10X, Y) = 10 \text{Cov}(X, Y)$.
- Pearson's Correlation standardizes the scale according to the standard deviations of X and Y :

$$\begin{aligned}\text{Corr}(X, Y) &= \mathbb{E} \left[\left(\frac{X - \mu_X}{\sigma_X} \right) \left(\frac{Y - \mu_Y}{\sigma_Y} \right) \right] \\ &= \frac{\text{Cov}(X, Y)}{\sqrt{\text{Var}(X) \text{Var}(Y)}}.\end{aligned}$$

- Can show that $-1 \leq \text{Corr}(X, Y) \leq 1$ using the Cauchy-Schwarz inequality
- a value of -1 implies a perfect negative linear relationship
- a value of 0 implies no linear relationship (not necessarily independent tho)
- a value of 1 implies a perfect linear relationship

Kendall's τ_K :

- Another measure of correlation
- Measures **monotonic** dependence instead of linear dependence
- Used on samples of observations, not on entire known distributions
- Measures concordance between each pair of observation (x_i, y_i) and (x_j, y_j) with $i \neq j$.
- Concordant means

$$\begin{aligned}x_i < x_j \quad \text{and} \quad y_i < y_j, \\ \text{or} \\ x_i > x_j \quad \text{and} \quad y_i > y_j;\end{aligned}$$

- Discordant means

$$\begin{aligned}x_i < x_j \quad \text{and} \quad y_i > y_j, \\ \text{or} \\ x_i > x_j \quad \text{and} \quad y_i < y_j;\end{aligned}$$

- The formal definition is then

$$\tau_K = \frac{\text{Number of concordant pairs} - \text{Number of discordant pairs}}{\binom{n}{2}},$$

where n is the sample size.

Mutual Information:

- Defined as

$$H(X, Y) = \sum_x \sum_y P(X = x \cap Y = y) \log \left[\frac{P(X = x \cap Y = y)}{P(X = x) \cdot P(Y = y)} \right].$$

- Not really used in D551.
- Apparently useful in Machine Learning.

Variance of a linear combination:

$$\text{Var}(aX + bY) = a^2 \text{Var}(X) + b^2 \text{Var}(Y) + 2ab \text{Cov}(X, Y).$$

Lecture 4

Conditional probabilities and conditional distributions

- In general, for events A and B , the conditional probability of A given B is

$$P(A | B) = \frac{P(A \cap B)}{P(B)}.$$

- If X and Y are two RVs, then $(X | Y = y)$ is a probability distribution

$$\sum_x P(X_{Y=y} = x) = \sum_x P(X = x | Y = y) = 1.$$

- And hence $Z = (X | Y = y)$ is also a random variable

Law of total expectation:

- A marginal mean can be computed from the conditional means and the probabilities of the conditioning variable.
- The formula, known as the **Law of Total Expectation**, is

$$\begin{aligned}\mathbb{E}_Y(Y) &= \sum_x \mathbb{E}_Y(Y | X = x) \cdot P(X = x) \\ &= \mathbb{E}_X[\mathbb{E}_Y(Y | X)].\end{aligned}$$

- Also note that the law of total probability implies that

$$\begin{aligned}P(Y = y) &= \sum_x P(Y = y | X = x) \cdot P(X = x) \\ &= \mathbb{E}_X[P(Y = y | X = x)].\end{aligned}$$

Conditional Independence:

- Note that even if X and Y are independent, they might not be independent when conditioning on some other RV Z .
- We say that X and Y are conditionally independent given Z if and only if

$$P(X = x \cap Y = y | Z = z) = P(X = x | Z = z) \cdot P(Y = y | Z = z) \quad \forall_{x,y,z}$$

- Equivalently:

$$P(X = x | Y = y, Z = z) = P(X = x | Z = z) \quad \forall_{x,y,z}$$

Example of variables that are not independent but are conditionally independent:

- See https://pages.github.ubc.ca/MDS-2022-23/DSCI_551_stat-prob-dsci_students/conditional-probabilities.html#conditional-independence
- Let L be a student's lab grade, Q be a student's quiz grade, and S represent whether the student is a Statistics major
- For simplicity, we will consider only Bernoulli random variables, so L and Q will only take on the values "high" and "low", and S takes on the values of "yes" or "no".
- L and Q appear to have a mild positive correlation: when the lab grade is high, the quiz grade is a bit more likely to also be high.

- So L and Q are not independent.
- But L and Q are conditionally independent given S !!
- Intuition: When L is high, it indicates that the student is more likely to be a statistics major, which in turn indicates that their quiz grade Q is more likely to be high. However, this is the **only** reason why high lab grades are correlated with high quiz grades. If you already know that a person is (or is not) a Statistics major, then their lab and quiz grades are completely independent.
- Note that it is also possible to have the opposite case: two variables that are marginally independent, but not conditionally independent given a third variable.

Lecture 5

Continuous random variables:

- Continuous random variables have an uncountably infinite number of possible outcomes
- In practice, we can never measure anything on a continuous scale since any measuring instrument must always round to some precision.
- But we can use a continuous distribution if the quantity being measured has “enough” precision that the distance between two neighbouring measurements is “not a big deal”.
- For example, if dealing with monetary quantities which are in the magnitude of millions of dollars, rounding to the nearest cent is no big deal.

Probability Density:

- Continuous RVs have an associated **probability density function** (PDF)
- It measures the density of probability per unit.
- The density of a RV X is denoted as $f_X(x)$
- In general,

$$P(a \leq X \leq b) = \int_a^b f_X(x) dx$$

- By the law of total probability, we have

$$\int_{-\infty}^{\infty} f_X(x) dx = 1.$$

- Note that

$$P(X = a) = \int_a^a f_X(x) dx = 0 \quad \forall_a$$

- A PDF must be non-negative everywhere

Distribution properties:

- See lecture 1 notes for definitions of mode, entropy, expectation, variance, and standard deviation.
- Recall that entropy can be negative for continuous RVs.

Median and Quantiles:

- The **median** is the outcome for which there is a 50-50 chance of seeing a greater (or lesser) value.
- By definition,

$$P[X \leq \text{Median}(X)] = 0.5 = P[X \geq \text{Median}(X)].$$

- Its empirical definition (i.e. the sample statistic version of median) is the “middle value” after sorting the observations from smallest to largest.
- Better than the mean in some ways because it is not as sensitive to outliers and reduces possibilities to two equally likely outcomes.
- The **p -quantile**, denoted $Q_X(p)$, is the outcome with a probability p of getting a smaller outcome.
- By definition,

$$P[X \leq Q_X(p)] = p.$$

- So the 0.5-quantile is just the median.

- An empirically-based definition of the p -quantile is the n th largest (rounded up) observation in a sample of size n .

Special Quantiles:

- The 0.25 , 0.5 , and 0.75 -quantiles are called **quartiles**.
- Often referred to as the 1st, 2nd, and 3rd quartiles, and denoted as Q_1, Q_2 , and Q_3 , respectively.
- The 0.01 , 0.02 , \dots , and 0.99 -quantiles are called **percentiles**.
- The p -quantile will often be called the 100th percentile; for example, the 40th percentile is the 0.4 -quantile, and the 97th percentile is the 0.97-quantile.

Prediction Intervals

- A $p \times 100\%$ prediction interval $[a, b]$ is such that

$$P(a \leq X \leq b) = p,$$

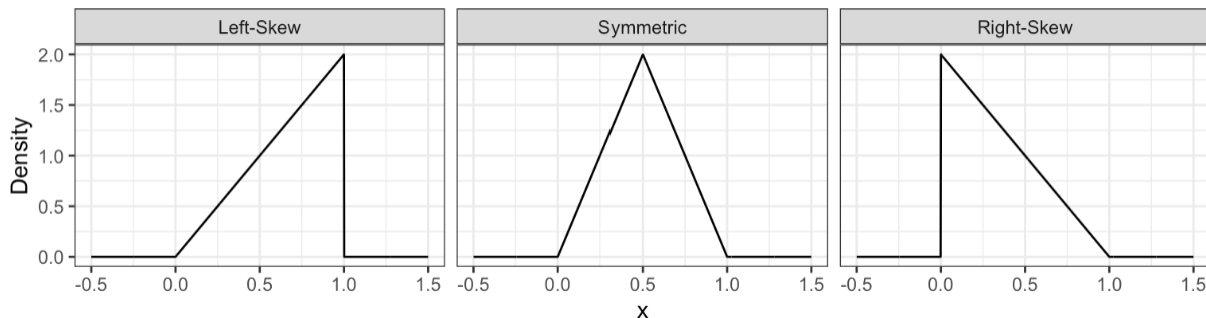
$$P(X \leq a) = \frac{1-p}{2},$$

$$P(X \geq b) = \frac{1-p}{2}.$$

- Clearly $a = Q_x\left(\frac{1-p}{2}\right)$
- Similarly, $b = Q_x\left(1 - \frac{1-p}{2}\right) = Q_x\left(\frac{1+p}{2}\right)$

Skewness

- **Skewness** measures how “lopsided” a distribution is, as well as the direction of the skew.
- If the density is symmetric, then the skewness is 0.
- If the density is more “spread-out” towards the right/positive values, then the distribution is said to be right-skewed (positive skewness).
- If the density is more “spread-out” towards the left/negative values, then the distribution is said to be left-skewed (negative skewness).



- The mathematical definition is given by

$$\text{Skewness}(X) = \mathbb{E} \left[\left(\frac{X - \mu_X}{\sigma_X} \right)^3 \right].$$

Representing Distributions:

- A PDF (or PMF) is just a representation of a distribution. It is not the distribution itself.
- There are alternative ways to represent a distribution.
- All of these representations capture everything about a distribution, meaning that if one of them is given, the other ones can be derived.

Cumulative Distribution Function

- The **cumulative distribution function** (CDF) is defined by

$$F_X(x) = P(X \leq x)$$

- It can be calculated as

$$F_X(x) = \int_{-\infty}^x f_X(t) dt.$$

- It is unitless (because probability is unitless)
- A valid CDF $F_X(x)$ must satisfy the following requirements:
 1. Must never decrease.
 2. It must never evaluate to be < 0 or > 1 .
 3. $F_X(x) \rightarrow 0$ as $x \rightarrow -\infty$
 4. $F_X(x) \rightarrow 1$ as $x \rightarrow \infty$.

Survival Function

- The **survival function** $S_X(x)$ is defined by

$$S_X(x) = P(X > x).$$

- It is the CDF “flipped upside down”:

$$S_X(x) = 1 - F_X(x)$$

Quantile Function

- The **quantile function** takes a probability p and maps it to the p -quantile.
- Mathematically, this means that

$$Q(p) = F^{-1}(p).$$

- The quantile function is the CDF with the axes swapped, or in other words, the CDF reflected diagonally.
- All these representations (PDF, CDF, survival function, quantile function) fully contain all the information about the distribution; they are different “views” of the same underlying random variable.

Lecture 6

Uniform Distribution:

$$X \sim \text{Uniform}(a, b).$$

- Has uniform density between a and b , where $a < b$. -PDF:

$$f_X(x | a, b) = \frac{1}{b - a} \quad \text{for } a \leq x \leq b.$$

Gaussian (Normal) Distribution:

$$X \sim \mathcal{N}(\mu, \sigma^2).$$

- μ is the mean, σ^2 is the variance
- Follows a “bell-shaped” curve
- PDF:

$$f_X(x | \mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left[-\frac{(x - \mu)^2}{2\sigma^2} \right] \quad \text{for } -\infty < x < \infty.$$

Log-Normal Distribution:

$$X \sim \text{Log-Normal}(\mu, \sigma^2).$$

- A random variable X is a Log-Normal distribution if the transformation $\ln(X)$ is Normal.
- Can't take the logarithm of a negative number, so must have $X \geq 0$.
- This means that

$$\ln(X) \sim \mathcal{N}(\mu, \sigma^2)$$

- PDF:

$$f_X(x | \mu, \sigma^2) = \frac{1}{x\sqrt{2\pi\sigma^2}} \exp \left\{ -\frac{[\log(x) - \mu]^2}{2\sigma^2} \right\} \quad \text{for } x \geq 0.$$

Exponential Distribution:

$$X \sim \text{Exponential}(\lambda).$$

- The exponential family is for positive random variables, often interpreted as wait time for some event to happen.
- Characterized by a memoryless property, where after waiting for a certain period of time, the remaining wait time has the same distribution.
- Parameterized by the average rate $\lambda > 0$ at which events happen.
- PDF:

$$f_X(x | \lambda) = \lambda e^{-\lambda x} \quad \text{for } x \geq 0.$$

Beta Distribution:

$$X \sim \text{Beta}(\alpha, \beta).$$

- The Beta family of distributions is defined for random variables taking values between 0 and 1, so is useful for modelling the distribution of proportions.
- Has the Uniform distribution as a special case.
- PDF:

$$f_X(x | \alpha, \beta) = \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} x^{\alpha-1} (1 - x)^{\beta-1} \quad \text{for } 0 \leq x \leq 1.$$

Weibull Distribution:

$$X \sim \text{Weibull}(\lambda, k).$$

- A generalization of the Exponential family, which allows for an event to be more likely the longer you wait.
- characterized by two parameters, a scale parameter $\lambda > 0$ and a shape parameter $k > 0$ (where $k = 1$ results in the Exponential family).
- PDF:

$$f_X(x | \lambda, k) = \frac{k}{\lambda} \left(\frac{x}{\lambda}\right)^{k-1} \exp^{-(x/\lambda)^k} \quad \text{for } x \geq 0.$$

Gamma Distribution:

$$X \sim \text{Gamma}(k, \theta).$$

- A two-parameter distribution family with support on non-negative numbers.
- PDF:

$$f_X(x | k, \theta) = \frac{1}{\Gamma(k)\theta^k} x^{k-1} \exp(-x/\theta) \quad \text{for } x \geq 0.$$

Relevant R Functions:

- Always take the form `<x><dist>`
- `<x>` is one of `d`, `p`, `q`, or `r`:
- `d`: density function (PDF), i.e. $f_X(x)$
- `p`: cumulative distribution function (CDF), i.e. $F_X(x)$
- `q`: quantile function (inverse of CDF), i.e. $Q_X(p)$
- `r`: random number generator.
- `<dist>` is an abbreviation of a distribution family
- Some abbreviations for `<dist>`:
- `unif`: Uniform (continuous).
- `norm`: Normal (continuous).
- `lnorm`: Log-Normal (continuous).
- `geom`: Geometric (discrete).
- `pois`: Poisson (discrete).
- `binom`: Binomial (discrete).

Continuous Conditional Distributions:

- In the continuous world, we replace probabilities with densities.
- For example, if X is continuous then

$$P(X = x | X \geq 2500) = \frac{P(X = x)}{P(X \geq 2500)}$$

is not a useful calculation, since $P(X = x) = 0$.

- Instead, we calculate the conditional density of X :

$$f_{X|X \geq 2500}(x) = \begin{cases} \frac{f_X(x)}{P(X \geq 2500)} & \text{for } x \geq 2500, \\ 0 & \text{for } x < 2500 \end{cases}.$$

- In general,

$$f_{X|X \geq a}(x) = \begin{cases} \frac{f_X(x)}{P(X \geq a)} & \text{for } x \geq a, \\ 0 & \text{for } x < a \end{cases}.$$

- Similarly, if Y is continuous then

$$P(X = x | Y = 5) = \frac{P(X = x)}{P(Y = 5)}$$

is not a useful calculation, since $P(Y = 5) = 0$.

- Instead, we calculate

$$f_{X|Y=5}(x) = \frac{f_{X,Y}(x, 5)}{f_Y(5)}.$$

- In general,

$$f_{X|Y}(x) = \frac{f_{X,Y}(x, y)}{f_Y(y)}.$$

Lecture 7

Continuous Independence:

- As per usual, in the continuous case we replace probabilities for densities
- Two continuous RVs X and Y are independent if

$$f_{X,Y}(x,y) = f_X(x) \cdot f_Y(y) \quad \forall_{x,y} \quad (1)$$

- Equivalently:

$$f_{Y|X}(y) = f_Y(y) \quad \forall_{x,y} \quad (2)$$

- Intuitively, (1) means that when slicing the joint density at various points along the x -axis or y -axis, the resulting 1-D function will be the same, up to some multiplication factor.
- e.g $f_{X,Y}(1,y) = f_X(1) \cdot f_Y(y)$ and $f_{X,Y}(2,y) = f_X(2) \cdot f_Y(y)$
- This means that all the slices along an axis will have the same shape
- Note that $f_{X,Y}(1,y)$ is **not** a PDF. The conditional PDF is

$$f_{Y|X=1}(y) = \frac{f_{X,Y}(1,y)}{f_X(1)} = f_Y(y),$$

which also follows directly from (2).

Direction of Dependence

- Contour plots are useful for determining if two RVs are dependent or not
- They also help determine the direction of dependence
- For two positively correlated variables, there is overall tendency of the contour lines to point up and to the right
- For two negatively correlated variables, there is overall tendency of the contour lines to point down and to the right

Estimating the marginal distribution from the conditional:

- By the law of total probability, we have

$$f_Y(y) = \int_x f_{X,Y}(x,y) \, dx = \int_x f_{Y|X}(y) f_X(x) \, dx.$$

- This is equivalent to

$$f_Y(y) = \mathbb{E}_X[f_{Y|X}(y)].$$

- If we know the conditional densities and have a sample of X as observed values x_1, \dots, x_n , then we can empirically estimate

$$f_Y(y) \approx \frac{1}{n} \sum_{i=1}^n f_{Y|X_i}(y).$$

- Similarly, for CDFs we have

$$F_Y(y) = \int_x F_{Y|X}(y) f_X(x) \, dx = \mathbb{E}_X[F_{Y|X}(y)].$$

- And we can estimate the CDF empirically with

$$F_Y(y) \approx \frac{1}{n} \sum_{i=1}^n F_{Y|X_i}(y).$$

Estimating the marginal mean from the conditional:

- As in the discrete case, we have

$$\mathbb{E}(Y) = \int_x m(x) f_X(x) dx = \mathbb{E}_X[\mathbb{E}_Y(Y | X = x)].$$

- Let $m(x) = \mathbb{E}(Y | X = x)$. Then the above becomes

$$\mathbb{E}(Y) = \mathbb{E}[m(X)]$$

- In machine learning, a function like $m(x)$ is called a **model function**.
- When you fit a model using Supervised Learning, you usually end up with an estimate of $m(x)$.
- Then, using the empirical mean, we have

$$E(Y) \approx \frac{1}{n} \sum_{i=1}^n m(x_i).$$

Estimating marginal quantiles from the conditional:

- Unfortunately, if you have the p -quantile of Y given $X = x$, then this does not give much info about the p -quantile of Y .
- This is because the average of quantiles is a horrible estimate for quantiles
- e.g. a mean of medians is a horrible estimate for the median.
- So we could calculate the p -quantile of Y as an average using the same ideas as above, but the estimate would be almost worthless.

Multivariate Gaussian Distribution:

- A multivariate distribution family in which:

1. All RVs must have Gaussian marginals
2. All linear combinations of the marginals are Gaussian (called **Gaussian dependence**)

- In general, it has two parameters:

1. A mean vector, $\vec{\mu}$
2. A covariance matrix, Σ

- $\vec{\mu}$ is just a vector of the marginal means.
- e.g. in the bivariate case we have

$$\vec{\mu} = \begin{pmatrix} \mu_X \\ \mu_Y \end{pmatrix}$$

- Σ is a symmetric matrix of marginal covariances, where the diagonal entries are the marginal variances.
- It is symmetric because $\text{Cov}(X, Y) = \text{Cov}(Y, X)$
- The diagonal entries are the marginal variances because $\text{Cov}(X, X) = \text{Var}(X)$.

- We can also look at a correlation matrix, \mathbf{P}
- Recall that

$$\rho_{XY} = \frac{\text{Cov}(X, Y)}{\sqrt{\text{Var}(X) \text{Var}(Y)}} = \frac{\sigma_{XY}}{\sqrt{\sigma_X^2 \sigma_Y^2}}.$$

- This implies that $\rho_{XX} = 1$.
- Thus in the bivariate case the correlation matrix is

$$\mathbf{P} = \begin{pmatrix} \rho_{XX} & \rho_{XY} \\ \rho_{XY} & \rho_{YY} \end{pmatrix} \\ = \begin{pmatrix} 1 & \rho_{XY} \\ \rho_{XY} & 1 \end{pmatrix}.$$

- Note that the covariance matrix is always used in the parameterization of a multivariate Gaussian distribution, although the correlation matrix has other uses (e.g PCA)
- The above means that to fully specify a d -dimensional Gaussian distribution, we need:
- the means and variances of all d random variables, and
- the $\binom{d}{2}$ covariance or correlations between each pair of random variables

Properties of the Multivariate Gaussian Distribution:

- For the multivariate Gaussian distribution, uncorrelated marginals implies independent marginals.
- Marginal distributions are Gaussian, by definition.
- The marginal distribution of a subset of variables can be obtained by just taking the relevant subset of means, and the relevant subset of the covariance matrix.
- Linear combinations of marginals are Gaussian, by definition.
- If we want to find the mean and variance of a linear combination, we apply the rules for expectation and variance of linear combinations:

$$\mathbb{E}[aX + bY] = a\mathbb{E}[X] + b\mathbb{E}[Y]$$

and

$$\text{Var}(aX + bY) = a^2 \text{Var}(X) + b^2 \text{Var}(Y) + 2ab \text{Cov}(X, Y).$$

- In the case of a multivariate Gaussian distribution this yields

$$\mathbb{E}[aX + bY] = a\mu_X + b\mu_Y$$

and

$$\text{Var}(aX + bY) = a^2\sigma_X^2 + b^2\sigma_Y^2 + 2ab\sigma_{XY}.$$

- In the case of the variance for linear combinations of $n \geq 3$ marginals, the formula will need to include all pairwise covariances.
- Conditional distributions are Gaussian.
- If (X, Y) have a bivariate Gaussian distribution then the distribution of Y given that $X = x$ is also Gaussian.

- Its mean is

$$\mu_{Y|X=x} = \mu_Y + \frac{\sigma_Y}{\sigma_X} \rho_{XY}(x - \mu_X).$$

- Its variance is

$$\sigma_{Y|X=x}^2 = (1 - \rho_{XY}^2) \sigma_Y^2.$$